

lung

July 23, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, \
    accuracy_score
```

```
[2]: df = pd.read_csv("/Users/hashithareddy/Desktop/Lung.csv")
```

```
[3]: df.head()
```

```
[3]:
```

	id	age	gender	country	diagnosis_date	cancer_stage	family_history	\
0	1	64.0	Male	Sweden	2016-04-05	Stage I	Yes	
1	2	50.0	Female	Netherlands	2023-04-20	Stage III	Yes	
2	3	65.0	Female	Hungary	2023-04-05	Stage III	Yes	
3	4	51.0	Female	Belgium	2016-02-05	Stage I	No	
4	5	37.0	Male	Luxembourg	2023-11-29	Stage I	No	

	smoking_status	bmi	cholesterol_level	hypertension	asthma	cirrhosis	\
0	Passive Smoker	29.4	199	0	0	1	
1	Passive Smoker	41.2	280	1	1	0	
2	Former Smoker	44.0	268	1	1	0	
3	Passive Smoker	43.0	241	1	1	0	
4	Passive Smoker	19.7	178	0	0	0	

	other_cancer	treatment_type	end_treatment_date	survived
0	0	Chemotherapy	2017-09-10	0
1	0	Surgery	2024-06-17	1
2	0	Combined	2024-04-09	0
3	0	Chemotherapy	2017-04-23	0
4	0	Combined	2025-01-08	0

```
[4]: df.info()  
df.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 890000 entries, 0 to 889999  
Data columns (total 17 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   id                    890000 non-null  int64  
1   age                  890000 non-null  float64  
2   gender               890000 non-null  object  
3   country              890000 non-null  object  
4   diagnosis_date       890000 non-null  object  
5   cancer_stage         890000 non-null  object  
6   family_history       890000 non-null  object  
7   smoking_status       890000 non-null  object  
8   bmi                  890000 non-null  float64  
9   cholesterol_level    890000 non-null  int64  
10  hypertension         890000 non-null  int64  
11  asthma               890000 non-null  int64  
12  cirrhosis            890000 non-null  int64  
13  other_cancer         890000 non-null  int64  
14  treatment_type       890000 non-null  object  
15  end_treatment_date   890000 non-null  object  
16  survived             890000 non-null  int64  
dtypes: float64(2), int64(7), object(8)  
memory usage: 115.4+ MB
```


```
[4]: id                    0  
age                      0  
gender                   0  
country                  0  
diagnosis_date           0  
cancer_stage             0  
family_history           0  
smoking_status           0  
bmi                      0  
cholesterol_level        0  
hypertension             0  
asthma                   0  
cirrhosis                0  
other_cancer             0  
treatment_type           0  
end_treatment_date       0  
survived                 0  
dtype: int64
```


```
[5]: df = df.drop(columns=["id", "diagnosis_date", "end_treatment_date"])

[6]: label_encoders = {}
for column in df.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le  # store for inverse_transform if needed later

[7]: # Correlation heatmap
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Feature Correlation")
plt.show()

# Target class distribution
sns.countplot(data=df, x='survived')
plt.title("Survival Distribution")
plt.show()
```





```
[8]: X = df.drop("survived", axis=1)
     y = df["survived"]

     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     ↪random_state=42)

[ ]: model = RandomForestClassifier(random_state=42)
     model.fit(X_train, y_train)

[ ]: y_pred = model.predict(X_test)

     print("Accuracy:", accuracy_score(y_test, y_pred))
     print("\nClassification Report:\n", classification_report(y_test, y_pred))

     # Confusion matrix
     sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blues')
     plt.xlabel("Predicted")
     plt.ylabel("Actual")
     plt.title("Confusion Matrix")
     plt.show()
```

```
[ ]: import joblib
      joblib.dump(model, "lung_cancer_survival_model.pkl")
```

```
[ ]: pip install xgboost
```

```
[ ]: from sklearn.linear_model import LogisticRegression
      from xgboost import XGBClassifier

      models = {
          "Logistic Regression": LogisticRegression(max_iter=1000),
          "Random Forest": RandomForestClassifier(random_state=42),
          "XGBoost": XGBClassifier(use_label_encoder=False, eval_metric='logloss',
          random_state=42)
      }

      for name, clf in models.items():
          clf.fit(X_train, y_train)
          y_pred = clf.predict(X_test)
          print(f"\n{name} Accuracy: {accuracy_score(y_test, y_pred):.4f}")
          print(classification_report(y_test, y_pred))
```

```
[ ]: # Feature Importance - Random Forest
      importances = model.feature_importances_
      features = X.columns

      # Sort and plot
      indices = np.argsort(importances)[::-1]
      plt.figure(figsize=(10,6))
      sns.barplot(x=importances[indices], y=features[indices])
      plt.title("Feature Importance - Random Forest")
      plt.show()
```

```
[35]: import joblib
      joblib.dump(model, "lung_cancer_survival_model.pkl")
```

```
[35]: ['lung_cancer_survival_model.pkl']
```

```
[ ]:
```